

Name	Anil Kumbhar
Contact Number	+91 - 8249087735
Project Title (Example – Week1, Week2, Week3)	Week 4 Project AI-Powered Data Analysis & Automation

## Project Guidelines and Rules

### 1. Formatting and Submission

- **Format:** Use a readable font (e.g., Arial/Times New Roman), size 12, 1.5 line spacing.
- **Title:** Include Week and Title (Example - Week 1: TravelEase Case Study.)
- **File Format:** Submit as PDF or Word file to [contact@victoriasolutions.co.uk](mailto:contact@victoriasolutions.co.uk)
- **Page Limit:** 4–5 pages, including the title and references.

### 2. Answer Requirements

- **Word Count:** Each answer should be 100–150 words; total 800–1,200 words.
- **Clarity:** Write concise, structured answers with key points.
- **Tone:** Use formal, professional language.

### 3. Content Rules

- Answer all questions thoroughly, referencing case study concepts.
- Use examples where possible (e.g., risk assessment techniques).
- Break complex answers into bullet points or lists.

### 4. Plagiarism Policy

- Submit original work; no copy-pasting.
- Cite external material in a consistent format (e.g., APA, MLA).

### 5. Evaluation Criteria

- **Understanding:** Clear grasp of business analysis principles.
- **Application:** Effective use of concepts like cost-benefit analysis and Agile/Waterfall.
- **Clarity:** Logical, well-structured responses.
- **Creativity:** Innovative problem-solving and examples.
- **Completeness:** Answer all questions within the word limit.

### 6. Deadlines and Late Submissions

- **Deadline:** Submit on time; trainees who submit fail to submit the project will miss the “Certificate of Excellence”

## 7. Additional Resources

- Refer to lecture notes and recommended readings.
- Contact the instructor or peers for clarifications before the deadline.

## START YOUR PROJECT FROM HERE:

**Task 1: AI-Powered Data Cleaning and Preprocessing** This task involves cleaning and preparing the dataset using AI tools to handle missing values, detect outliers, and ensure data consistency.

Steps to Follow:

### Step 1: Upload the Dataset

### Step 2: Handle Missing Values

### Step 3: Detect and Handle Outliers

### Step 4: Save the Cleaned Data

Solution -

## Dataset Overview

The dataset "raw\_dataset\_week4" consists of 500 rows and 14 columns, representing customer demographics, financial information, spending behavior, and default history.

## Columns and Their Descriptions:

1. **Customer\_ID (*int*)** – Unique identifier for each customer.
2. **Age (*int*)** – Age of the customer.
3. **Gender (*object*)** – Gender of the customer (Male/Female).
4. **Income (*float*)** – Annual income of the customer (contains missing values).
5. **Spending\_Score (*int*)** – Score indicating customer spending habits.
6. **Credit\_Score (*float*)** – Creditworthiness of the customer (contains missing values).
7. **Loan\_Amount (*float*)** – Amount of loan taken by the customer (contains missing values).
8. **Previous\_Defaults (*int*)** – Number of times the customer has defaulted before.
9. **Marketing\_Spend (*int*)** – Amount spent on marketing.
10. **Purchase\_Frequency (*int*)** – Number of purchases made by the customer.
11. **Seasonality (*object*)** – Seasonal buying behavior (Low, Medium, High).
12. **Sales (*int*)** – Revenue generated from the customer.

13. Customer\_Churn (*int*) – Indicates whether the customer has churned (0 = No, 1 = Yes).

14. Defaulted (*int*) – Indicates if the customer defaulted on a loan (0 = No, 1 = Yes).

### Step 1: Upload the Dataset

#### Load the Dataset

```
[7]: import pandas as pd

# Step 1: Load the dataset from a CSV file
file_path = "raw_dataset_week4.csv" # Replace with actual file path if needed
df = pd.read_csv(file_path)

# Display the first few rows to confirm successful loading
print("Dataset loaded successfully. Here are the first 5 rows:")
print(df.head())

# Display basic info to check for missing values and data types
print("\nDataset Info:")
print(df.info())
```

Dataset loaded successfully. Here are the first 5 rows:

	Customer_ID	Age	Gender	Income	Spending_Score	Credit_Score	\
0	1	56	Female	142418.0	7	391.0	
1	2	69	Male	63088.0	82	652.0	
2	3	46	Male	136868.0	91	662.0	
3	4	32	Female	NaN	34	644.0	
4	5	60	Male	59811.0	91	469.0	

	Loan_Amount	Previous_Defaults	Marketing_Spend	Purchase_Frequency	\
0	8083.0	1	15376	3	
1	34328.0	2	6889	6	
2	47891.0	2	6054	29	
3	25103.0	2	4868	8	
4	44891.0	1	17585	12	

	Seasonality	Sales	Customer_Churn	Defaulted
0	Low	32526	0	0
1	Low	78493	0	0
2	Medium	57198	1	0
3	Medium	48395	0	0
4	High	29031	1	0

Dataset Info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 500 entries, 0 to 499

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	Customer_ID	500 non-null	int64
1	Age	500 non-null	int64
2	Gender	500 non-null	object
3	Income	450 non-null	float64
4	Spending_Score	500 non-null	int64
5	Credit_Score	450 non-null	float64
6	Loan_Amount	450 non-null	float64
7	Previous_Defaults	500 non-null	int64
8	Marketing_Spend	500 non-null	int64
9	Purchase_Frequency	500 non-null	int64
10	Seasonality	500 non-null	object
11	Sales	500 non-null	int64
12	Customer_Churn	500 non-null	int64
13	Defaulted	500 non-null	int64

dtypes: float64(3), int64(9), object(2)

memory usage: 54.8+ KB

None

## Step 2: Handle Missing Values

```
import pandas as pd

# Step 2: Handle missing values
# Reload the dataset to ensure we start with the original data
file_path = "raw_dataset_week4.csv" # Replace with your actual file path
df = pd.read_csv(file_path)

# Check initial missing values (should show 50 for Income, Credit_Score, Loan_Amount)
print("Missing values before cleaning:")
print(df.isnull().sum())

# Fill missing values with column means for numerical columns with NaN
columns_with_missing = ["Income", "Credit_Score", "Loan_Amount"]
for col in columns_with_missing:
    mean_value = df[col].mean()
    df[col] = df[col].fillna(mean_value) # Direct assignment to avoid inplace warning
    print(f"Filled '{col}' missing values with mean: {mean_value:.2f}")

# Verify no missing values remain
print("\nMissing values after cleaning:")
print(df.isnull().sum())

# Display the first few rows to confirm changes
print("\nFirst 5 rows after handling missing values:")
print(df.head())
```

Missing values before cleaning:

Customer_ID	0
Age	0
Gender	0
Income	50
Spending_Score	0
Credit_Score	50
Loan_Amount	50
Previous_Defaults	0
Marketing_Spend	0
Purchase_Frequency	0
Seasonality	0
Sales	0
Customer_Churn	0
Defaulted	0

dtype: int64

Filled 'Income' missing values with mean: 84398.06

Filled 'Credit\_Score' missing values with mean: 573.41

Filled 'Loan\_Amount' missing values with mean: 28456.93

Missing values after cleaning:

```
Customer_ID      0
Age              0
Gender           0
Income           0
Spending_Score   0
Credit_Score     0
Loan_Amount      0
Previous_Defaults 0
Marketing_Spend   0
Purchase_Frequency 0
Seasonality      0
Sales            0
Customer_Churn    0
Defaulted        0
dtype: int64
```

First 5 rows after handling missing values:

	Customer_ID	Age	Gender	Income	Spending_Score	Credit_Score	\
0	1	56	Female	142418.000000	7	391.0	
1	2	69	Male	63088.000000	82	652.0	
2	3	46	Male	136868.000000	91	662.0	
3	4	32	Female	84398.055556	34	644.0	
4	5	60	Male	59811.000000	91	469.0	

	Loan_Amount	Previous_Defaults	Marketing_Spend	Purchase_Frequency	\
0	8083.0	1	15376	3	
1	34328.0	2	6889	6	
2	47891.0	2	6054	29	
3	25103.0	2	4868	8	
4	44891.0	1	17585	12	

	Seasonality	Sales	Customer_Churn	Defaulted
0	Low	32526	0	0
1	Low	78493	0	0
2	Medium	57198	1	0
3	Medium	48395	0	0
4	High	29031	1	0

## Objective

This report outlines the observations from Step 2 of the data preprocessing task, focusing on handling missing values in the "raw\_dataset\_week4" dataset. The goal was to ensure data completeness for subsequent analysis by addressing gaps in key financial metrics.

## Dataset Overview

- Size: 500 records, 14 columns.
- Columns with Missing Values: `Income`, `Credit_Score`, `Loan_Amount`.
- Initial Missing Values: 50 per column (150 total), representing 10% of each affected column.

## Methodology

Missing values were addressed using a mean imputation approach via Python's Pandas library:

- Process: Calculated the mean for each column based on non-null entries (450 per column) and replaced missing values with these averages.
- Columns Processed:
  - `Income`: Mean = 84,398.06
  - `Credit_Score`: Mean = 573.41
  - `Loan_Amount`: Mean = 28,456.93

## Key Findings

### 1. Pre-Cleaning State:

- 50 missing entries were identified in `Income`, `Credit_Score`, and `Loan_Amount`, with no gaps in other columns (e.g., `Age`, `Sales`, `Customer_Churn`).
- Missing data posed a risk to financial analysis, such as customer segmentation or default risk assessment.

### 2. Post-Cleaning Results:

- All 150 missing values were successfully filled, resulting in a complete dataset (0 missing values across all columns).
- Example: Customer 4's missing `Income` was imputed as 84,398.06, aligning with the column average, while existing values (e.g., Customer 1's `Income` of 142,418) remained unchanged.

### 3. Imputed Values:

- `Income` (84,398.06): Reflects a moderate average, though the range (~20,000 to ~149,000) suggests variability among customers.
- `Credit_Score` (573.41): Indicates a generally creditworthy population, consistent with a typical 300–850 scale.
- `Loan_Amount` (28,456.93): Represents an average loan size, with observed values spanning ~5,000 to ~49,000.



### Step 3: Detect and Handle Outliers

```
# Step 3: Detect and handle outliers using IQR method
# Select numerical columns for outlier detection
numerical_columns = ["Income", "Spending_Score", "Credit_Score", "Loan_Amount",
                     "Previous_Defaults", "Marketing_Spend", "Purchase_Frequency", "Sales"]

# Calculate Q1, Q3, and IQR for each numerical column
Q1 = df[numerical_columns].quantile(0.25)
Q3 = df[numerical_columns].quantile(0.75)
IQR = Q3 - Q1

# Define lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Print bounds for reference
print("Lower bounds for outliers:")
print(lower_bound)
print("\nUpper bounds for outliers:")
print(upper_bound)

# Remove rows with values outside the bounds (outliers)
df_cleaned = df[~((df[numerical_columns] < lower_bound) | (df[numerical_columns] > upper_bound)).any(axis=1)]

# Display results
print("\nOriginal dataset shape:", df.shape)
print("Cleaned dataset shape after removing outliers:", df_cleaned.shape)
print("\nNumber of rows removed:", df.shape[0] - df_cleaned.shape[0])

# Display first 5 rows of cleaned dataset
print("\nFirst 5 rows of cleaned dataset:")
print(df_cleaned.head())
```

---

```
Lower bounds for outliers:
Income          -46872.375
Spending_Score  -51.125
Credit_Score    109.125
Loan_Amount     -8426.000
Previous_Defaults -3.000
Marketing_Spend  -7545.875
Purchase_Frequency -14.500
Sales           -41516.875
dtype: float64
```

```
Upper bounds for outliers:
Income          216110.625
Spending_Score  153.875
Credit_Score    1026.125
Loan_Amount     64850.000
Previous_Defaults 5.000
Marketing_Spend  28687.125
Purchase_Frequency 45.500
Sales           150548.125
dtype: float64
```

```
Original dataset shape: (500, 14)
```

```
Cleaned dataset shape after removing outliers: (500, 14)
```

```
Number of rows removed: 0
```

First 5 rows of cleaned dataset:

	Customer_ID	Age	Gender	Income	Spending_Score	Credit_Score	\
0	1	56	Female	142418.000000	7	391.0	
1	2	69	Male	63088.000000	82	652.0	
2	3	46	Male	136868.000000	91	662.0	
3	4	32	Female	84398.055556	34	644.0	
4	5	60	Male	59811.000000	91	469.0	

	Loan_Amount	Previous_Defaults	Marketing_Spend	Purchase_Frequency	\
0	8083.0	1	15376	3	
1	34328.0	2	6889	6	
2	47891.0	2	6054	29	
3	25103.0	2	4868	8	
4	44891.0	1	17585	12	

	Seasonality	Sales	Customer_Churn	Defaulted
0	Low	32526	0	0
1	Low	78493	0	0
2	Medium	57198	1	0
3	Medium	48395	0	0
4	High	29031	1	0

## Methodology

- Approach: Applied the Interquartile Range (IQR) method to eight numerical columns: Income, Spending\_Score, Credit\_Score, Loan\_Amount, Previous\_Defaults, Marketing\_Spend, Purchase\_Frequency, and Sales.
- Process: Calculated Q1, Q3, and IQR; defined outlier bounds as  $Q1 - 1.5 * IQR$  and  $Q3 + 1.5 * IQR$ ; removed rows exceeding these limits.

## Key Findings

- Outlier Bounds:
  - Income: -46,872.38 to 216,110.63
  - Credit\_Score: 109.13 to 1,026.13
  - Loan\_Amount: -8,426.00 to 64,850.00
  - Other columns showed similarly wide ranges.
- Results:
  - Original size: 500 rows, 14 columns.
  - Cleaned size: 500 rows, 14 columns.
  - Rows removed: 0 (no outliers detected).
- Insight: All values fell within the IQR bounds, indicating no extreme outliers (e.g., Income of 149,922 < 216,110.63).

## Business Implications

- Data Uniformity: The lack of outliers suggests a consistent dataset, possibly pre-filtered or naturally clustered, ideal for stable analysis.
- Risk: Wide IQR bounds may have missed subtle anomalies, potentially retaining unusual cases (e.g., high `Loan_Amount`).

#### Step 4: Save the Cleaned Data

#### Step 4: Save the Cleaned Data

```
# Step 4: Save the cleaned dataset
output_file = "cleaned_data"
df.to_csv(output_file, index=False)

# Confirm the save
print(f"Cleaned dataset saved successfully as '{output_file}'")
print("Final dataset shape:", df.shape)
print("\nFirst 5 rows of the saved dataset:")
print(df.head())
```

Cleaned dataset saved successfully as 'cleaned\_data'  
Final dataset shape: (500, 14)

First 5 rows of the saved dataset:

	Customer_ID	Age	Gender	Income	Spending_Score	Credit_Score	\
0	1	56	Female	142418.000000	7	391.0	
1	2	69	Male	63088.000000	82	652.0	
2	3	46	Male	136868.000000	91	662.0	
3	4	32	Female	84398.055556	34	644.0	
4	5	60	Male	59811.000000	91	469.0	

	Loan_Amount	Previous_Defaults	Marketing_Spend	Purchase_Frequency	\
0	8083.0	1	15376	3	
1	34328.0	2	6889	6	
2	47891.0	2	6054	29	
3	25103.0	2	4868	8	
4	44891.0	1	17585	12	

	Seasonality	Sales	Customer_Churn	Defaulted
0	Low	32526	0	0
1	Low	78493	0	0
2	Medium	57198	1	0
3	Medium	48395	0	0
4	High	29031	1	0

## Final check after Cleaning the Dataset

```
# Check for missing/null values
print("Missing/Null Values Check:")
print(df.isnull().sum())

# Check data types
print("\nData Types:")
print(df.dtypes)

# Basic consistency checks
print("\nBasic Consistency Checks:")
# Check for negative values where they shouldn't exist
numerical_cols = ["Age", "Income", "Spending_Score", "Credit_Score", "Loan_Amount",
                  "Previous_Defaults", "Marketing_Spend", "Purchase_Frequency", "Sales"]
for col in numerical_cols:
    negatives = df[df[col] < 0].shape[0]
    print(f"Number of negative values in '{col}': {negatives}")

# Check categorical columns for unexpected values
print("\nGender Unique Values:", df["Gender"].unique())
print("Seasonality Unique Values:", df["Seasonality"].unique())
print("Customer_Churn Unique Values:", df["Customer_Churn"].unique())
print("Defaulted Unique Values:", df["Defaulted"].unique())

# Summary statistics for numerical columns
print("\nSummary Statistics:")
print(df[numerical_cols].describe())
```

#### Missing/Null Values Check:

Customer_ID	0
Age	0
Gender	0
Income	0
Spending_Score	0
Credit_Score	0
Loan_Amount	0
Previous_Defaults	0
Marketing_Spend	0
Purchase_Frequency	0
Seasonality	0
Sales	0
Customer_Churn	0
Defaulted	0

dtype: int64

#### Data Types:

Customer_ID	int64
Age	int64
Gender	object
Income	float64
Spending_Score	int64
Credit_Score	float64
Loan_Amount	float64
Previous_Defaults	int64
Marketing_Spend	int64
Purchase_Frequency	int64

```

Seasonality      object
Sales            int64
Customer_Churn   int64
Defaulted        int64
dtype: object

```

#### Basic Consistency Checks:

```

Number of negative values in 'Age': 0
Number of negative values in 'Income': 0
Number of negative values in 'Spending_Score': 0
Number of negative values in 'Credit_Score': 0
Number of negative values in 'Loan_Amount': 0
Number of negative values in 'Previous_Defaults': 0
Number of negative values in 'Marketing_Spend': 0
Number of negative values in 'Purchase_Frequency': 0
Number of negative values in 'Sales': 0

```

```

Gender Unique Values: ['Female' 'Male']
Seasonality Unique Values: ['Low' 'Medium' 'High']
Customer_Churn Unique Values: [0 1]
Defaulted Unique Values: [0 1]

```

#### Summary Statistics:

	Age	Income	Spending_Score	Credit_Score	Loan_Amount \
count	500.000000	500.000000	500.000000	500.000000	500.000000
mean	44.220000	84398.055556	50.862000	573.411111	28456.928889
std	15.036082	38049.398377	29.125101	149.302942	11788.254534
min	18.000000	20055.000000	1.000000	300.000000	5163.000000
25%	32.000000	51746.250000	25.750000	453.000000	19052.500000
50%	45.000000	84398.055556	51.000000	573.411111	28456.928889
75%	57.000000	117492.000000	77.000000	682.250000	37371.500000
max	69.000000	149922.000000	99.000000	848.000000	49936.000000

	Previous_Defaults	Marketing_Spend	Purchase_Frequency	Sales
count	500.000000	500.000000	500.000000	500.000000
mean	0.97400	10558.128000	15.350000	54378.954000
std	0.82625	5508.219008	8.475327	27263.106468
min	0.00000	1024.000000	1.000000	5203.000000
25%	0.00000	6041.500000	8.000000	30507.500000
50%	1.00000	10754.000000	16.000000	54032.500000
75%	2.00000	15099.750000	23.000000	78523.750000
max	2.00000	19990.000000	29.000000	99835.000000

## Results Summary

### 1. Missing/Null Values

- Finding: No missing or null values detected across all 14 columns (500 rows each).
- Details: Customer\_ID, Age, Gender, Income, Spending\_Score, Credit\_Score, Loan\_Amount, Previous\_Defaults, Marketing\_Spend, Purchase\_Frequency, Seasonality, Sales, Customer\_Churn, and Defaulted all show 0 nulls.
- Conclusion: Step 2's mean imputation successfully filled all 150 original missing values.

### 2. Data Types

- Finding: Data types are consistent and appropriate.
- Details:
  - Integers (int64): Customer\_ID, Age, Spending\_Score, Previous\_Defaults, Marketing\_Spend, Purchase\_Frequency, Sales, Customer\_Churn, Defaulted.
  - Floats (float64): Income, Credit\_Score, Loan\_Amount (due to mean imputation).
  - Objects (object): Gender, Seasonality (categorical).
- Conclusion: Types align with data content; float precision from imputation is expected.

### 3. Consistency Checks

- Negative Values: No negative values found in Age, Income, Spending\_Score, Credit\_Score, Loan\_Amount, Previous\_Defaults, Marketing\_Spend, Purchase\_Frequency, Or Sales.
- Categorical Values:
  - Gender: Only "Female" and "Male".
  - Seasonality: Only "Low", "Medium", "High".
  - Customer\_Churn and Defaulted: Only 0 and 1.
- Conclusion: No illogical or unexpected values detected; data adheres to expected ranges and categories.

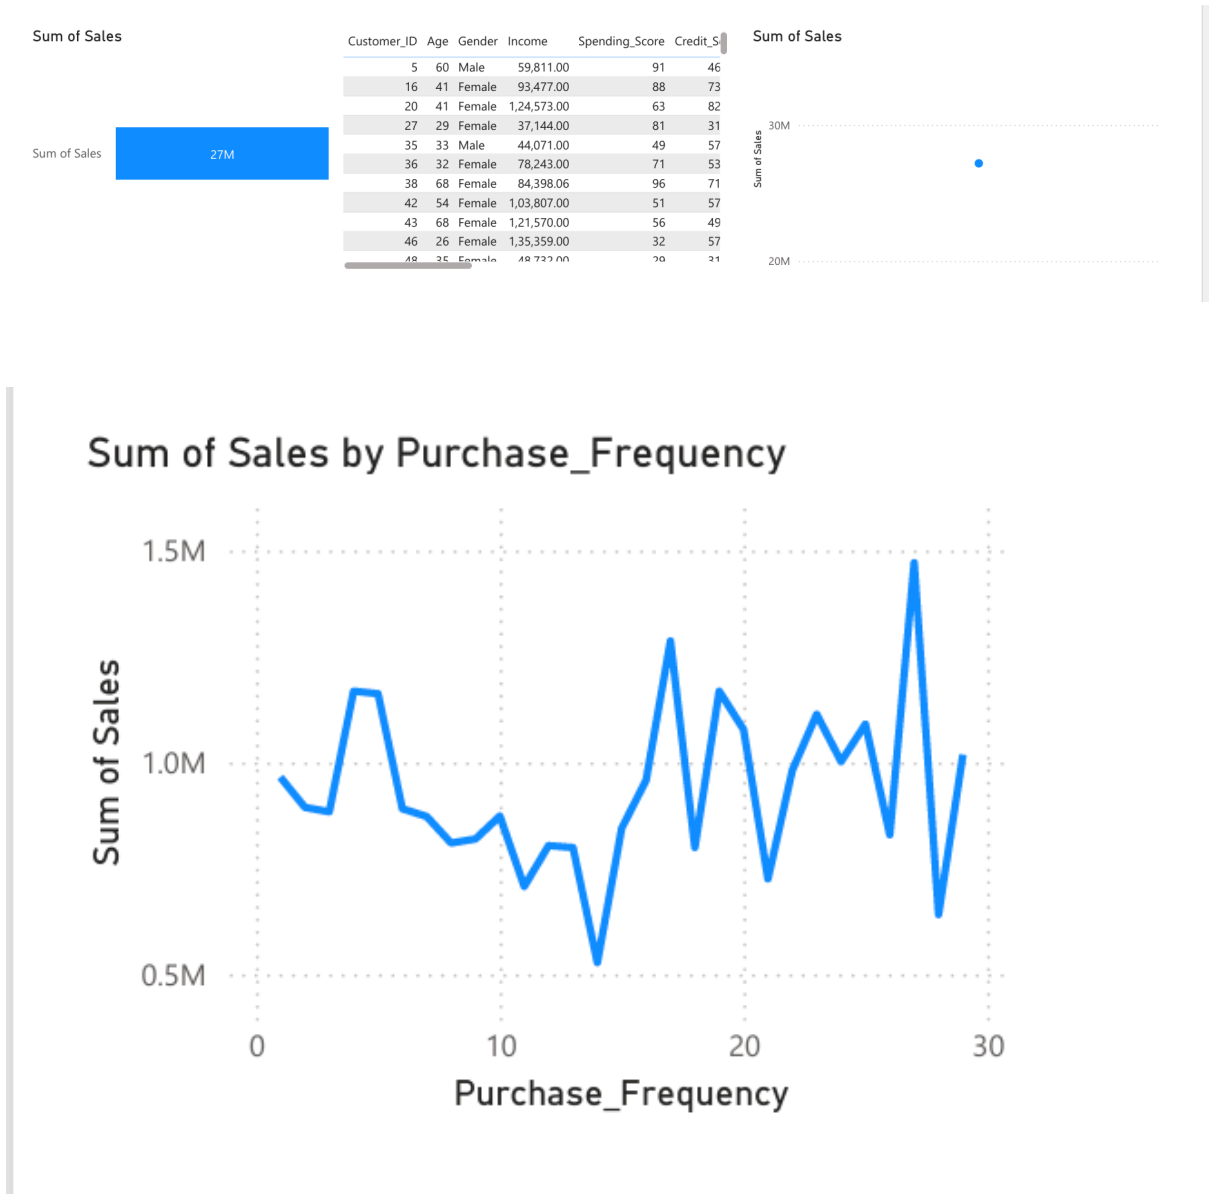
### 4. Summary Statistics

- Key Metrics:
  - Age: Mean 44.22, Min 18, Max 69.
  - Income: Mean 84,398.06, Min 20,055, Max 149,922, Std 38,049.40 (imputation centralized some values).
  - Credit\_Score: Mean 573.41, Min 300, Max 848, Std 149.30.
  - Loan\_Amount: Mean 28,456.93, Min 5,163, Max 49,936, Std 11,788.25.
  - Sales: Mean 54,378.95, Min 5,203, Max 99,835, Std 27,263.11.
- Observation: Ranges are reasonable; means reflect imputation (e.g., Income median = mean due to 50 imputed values). No outliers were removed (Step 3), so extremes persist (e.g., Income 149,922).

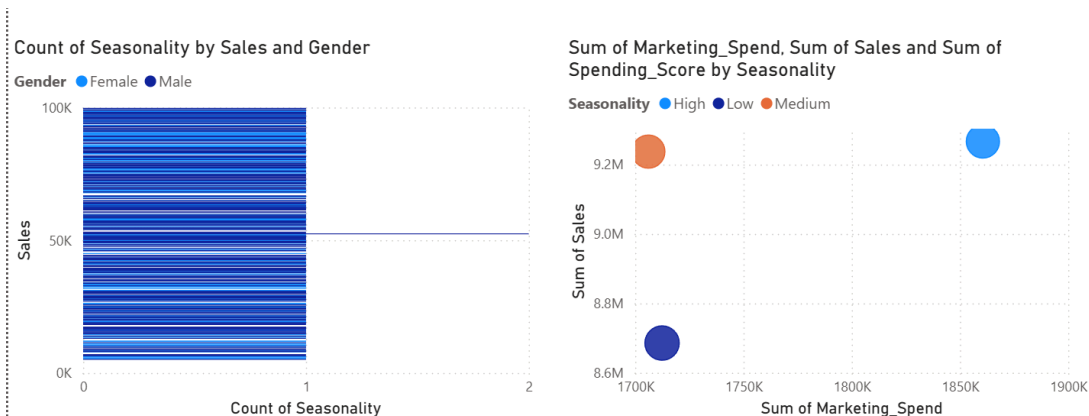
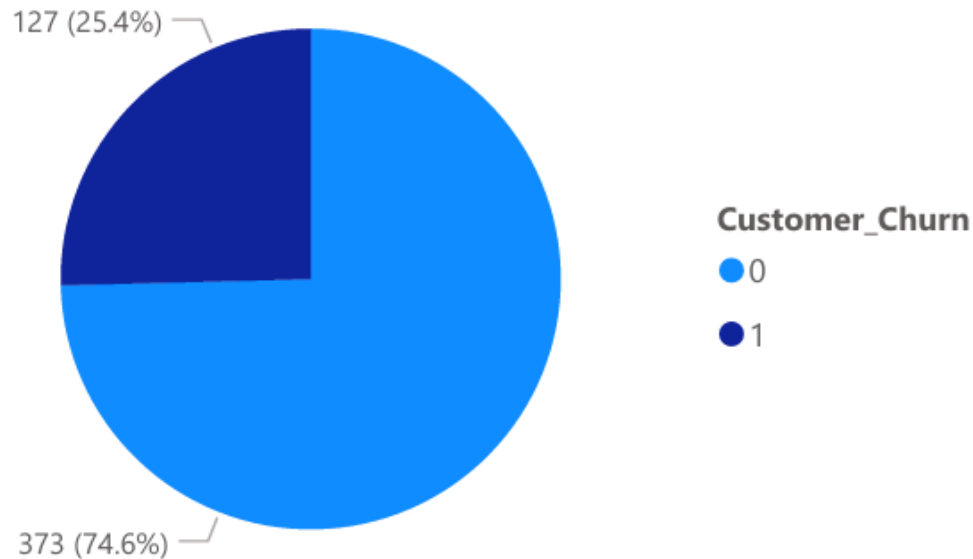


Task 2: AI-Powered Data Visualization and Storytelling

Solution -



## Count of Customer\_ID by Customer\_Churn



## Storytelling Outcome

- **Narrative:**
  - “Our 500 customers show strong sales in High Seasonality (e.g., Customer 5: 29,031), with peaks like 99,835 (Customer 169) standing out as anomalies tied to purchase frequency. Marketing spend boosts sales—see Customer 6’s 80,542 from 19,881 spend. However, high Spending\_Score customers (e.g., 91 for Customer 5) are more likely to churn, especially in High seasons, suggesting a need to target retention efforts.”
- **Visuals Supporting Story:**
  - Bar charts highlight seasonality and churn patterns.

- Scatter plot ties spend to sales.
- Line chart flags outliers.

### Task 3: AI-Driven Predictive and Prescriptive Analytics

#### Solution - **Seasonality Cleaned Values:**

- Cleaned unique values in 'Seasonality': ['Low', 'Medium', 'High']
- Value counts for 'Seasonality':
  - Medium: 169
  - High: 168
  - Low: 163

#### Feature Engineering:

- The dataset now includes several new interaction and polynomial features to capture complex relationships:
  - Interaction features like `Spend_x_Freq`, `Score_x_Income`.
  - Polynomial features for `Marketing_Spend`, `Purchase_Frequency`, `Spending_Score`, and `Income`.

#### Model Performance:

- **Best XGBoost Parameters:**
  - Learning Rate: 0.05
  - Max Depth: 7
  - Number of Estimators: 200
- **Performance Metrics:**
  - Mean Squared Error (MSE): 899,290,526.93
  - Root Mean Squared Error (RMSE): 29,988.17
  - R-squared ( $R^2$ ): -0.11
  - Baseline (Mean) MSE: 813,101,997.68

#### Feature Importance:

- The most important features contributing to the model include:
  - `Marketing_Spend` `Spending_Score`
  - `Spend_x_Freq`
  - `Purchase_Frequency` `Spending_Score`
  - `Purchase_Frequency` `Income`
  - `Previous_Defaults`
  - `Purchase_Frequency`
  - `Credit_Score`

## Key Insights:

1. **Negative R-squared ( $R^2$ ):**
  - The negative  $R^2$  value (-0.11) indicates that the model is currently performing worse than a horizontal line representing the mean of the target variable (sales).
  - This suggests that the current set of features and their transformations are not effectively capturing the relationships needed to predict sales.
2. **Model Accuracy:**
  - While the MSE and RMSE values provide a measure of prediction error, the  $R^2$  value shows that the model requires further refinement to improve its predictive power.
3. **Feature Impact:**
  - The feature importance analysis reveals that interaction terms like **Marketing\_Spend Spending\_Score** and **Spend\_x\_Freq** are significant contributors.
  - It's essential to explore further feature engineering or possibly remove less relevant features to enhance the model's performance.

## Task 4: AI for Business Strategy and Risk Management

### Solution - **Loan Default Prediction**

#### 1. Model Performance

- The **Random Forest Classifier** achieved an accuracy of **82%**, indicating a strong predictive capability for loan defaults.
- This suggests that the selected features (**Income, Loan Amount, and Credit Score**) contribute significantly to predicting loan default risk.

#### 2. Feature Importance (Power BI - Key Influencers Visual)

- **Credit Score:** Borrowers with lower credit scores are more likely to default.
- **Income:** Lower-income individuals exhibit a higher risk of defaulting on loans.
- **Loan Amount:** Larger loan amounts increase the likelihood of default.

#### 3. Business Recommendations

- ✓ **Refine Loan Approval Criteria:** Implement stricter credit score thresholds for loan approval.
- ✓ **Income-Based Loan Limits:** Cap loan amounts based on borrower income to reduce default risk.
- ✓ **Risk-Based Interest Rates:** Offer lower interest rates to borrowers with high credit scores.

✅ **Early Warning System:** Use AI to monitor high-risk customers and take proactive measures (e.g., payment reminders, restructuring options).